

①

$$Job = 1$$

$$CPU = 1$$

$$\text{run-time} = 30$$

$$\text{working-set size} = 200$$

$$-n \mid -L \ a: 30 : 200$$

Working-set size  $\Rightarrow$  how much cache space it needs to run efficiently

$$\text{Finished time} = 30$$

$$CPU \cdot \text{utilization} = 100$$

warm = 0.00  $\Rightarrow$  since default cache size is 100

with cache size 200

$$-n \mid -M 200 \ -L \ a: 30 : 200$$

We get -

Finished time = 20  $\Rightarrow$  since warm-rate

is 2x, now it runs faster on CPU  
with a warm cache

My  
experience  
not part  
of sim

*the next*

CPU utilization = 100

$$\text{warm} = \frac{10}{20} = \frac{1}{2} = 50\%$$

warm-time is 10, hence after 10.  
 Remaining 20, are done in 10  
 due to warm-rate 2x

②

-n 1 -L a; 30:200 -M 300  
 Did already above

③

-n 1 -M 300 -L a; 30:200 -T  
 Reduces 1x then 2x due to  
 warm-rate for next 10 seconds

Scheduler central queue: ['a']		
0	a	[ 29]
1	a	[ 28]
2	a	[ 27]
3	a	[ 26]
4	a	[ 25]
5	a	[ 24]
6	a	[ 23]
7	a	[ 22]
8	a	[ 21]
9	a	[ 20]
<hr/>		
10	a	[ 18]
11	a	[ 16]
12	a	[ 14]
13	a	[ 12]
14	a	[ 10]
15	a	[ 8]
16	a	[ 6]
17	a	[ 4]
18	a	[ 2]
19	a	[ 0]

-n 1 -M 300 -L a:30:200 -r 4-T  
Now reduces with 4x  $\frac{1}{4}$   
warm-rate

```
Scheduler central queue: ['a']  
0 a [ 29]  
1 a [ 28]  
2 a [ 27]  
3 a [ 26]  
4 a [ 25]  
5 a [ 24]  
6 a [ 23]  
7 a [ 22]  
8 a [ 21]  
9 a [ 20]  
-----  
10 a [ 16]  
11 a [ 12]  
12 a [ 8]  
13 a [ 4]  
14 a [ 0]
```

④ -n 1 -M 300 -L a:30:200 -C

```
0 a cache[ ]  
1 a cache[ ]  
2 a cache[ ]  
3 a cache[ ]  
4 a cache[ ]  
5 a cache[ ]  
6 a cache[ ]  
7 a cache[ ]  
8 a cache[ ]  
9 a cache[w]  
-----  
10 a cache[w]  
11 a cache[w]  
12 a cache[w]  
13 a cache[w]  
14 a cache[w]  
15 a cache[w]  
16 a cache[w]  
17 a cache[w]  
18 a cache[w]  
19 a cache[w]
```

Default warm-time = 10

-n 1 -M 300 -L a:30:200 -r 5-T  $\downarrow$

warm-time

```
0  a cache[ ]  
1  a cache[ ]  
2  a cache[ ]  
3  a cache[ ]  
4  a cache[w]  
5  a cache[w]  
6  a cache[w]  
7  a cache[w]  
8  a cache[w]  
9  a cache[w]  
  
-----  
10 a cache[w]  
11 a cache[w]  
12 a cache[w]  
13 a cache[w]  
14 a cache[w]  
15 a cache[w]  
16 a cache[w]  
17 a cache[w]
```

-n | -M 300 -L a:30:200 -w 15 -c

Time  $\Rightarrow$  15 (to warm ~~is~~, 15 remaining  
cache)

with warm-rate = 2

$$15 + \frac{15}{2} = 15 + 8 = \underline{\underline{23}}$$

```
0  a cache[ ]  
1  a cache[ ]  
2  a cache[ ]  
3  a cache[ ]  
4  a cache[ ]  
5  a cache[ ]  
6  a cache[ ]  
7  a cache[ ]  
8  a cache[ ]  
9  a cache[ ]  
  
-----  
10 a cache[ ]  
11 a cache[w]  
12 a cache[w]  
13 a cache[w]  
14 a cache[w]  
15 a cache[w]  
16 a cache[w]  
17 a cache[w]  
18 a cache[w]  
19 a cache[w]  
  
-----  
20 a cache[w]  
21 a cache[w]  
22 a cache[w]
```

⑤ -n 2 -L a:100:100,b:100:50,c:100:50  
 ↓  
 CPUs 3 jobs, cache-size = 100 (default)  
 quantum = 10 Cache space needed to run efficiently  
 Job a  $\Rightarrow$  runtime = 100, working-set-size = 100

$$\text{Job b} \Rightarrow \begin{matrix} 1 \\ | \\ 1 \end{matrix} = 50$$

$$\text{Job c} \Rightarrow \begin{matrix} 1 \\ | \\ 1 \end{matrix} = 50 \quad \text{warm-time} = 10$$

Round robin in seq. for multiple CPUs

CPU 1	CPU 2
Job a 0-9 cache warmed for Job a = 90 Job a	Job b 0-9 cache warmed for Job b = 90 Job b

Job c  
Job b  
Job a

Job a  
Job c  
Job b

No chance of warm cache since cache-size is 100

Other job run overwrites that cache

Total = 300  $\Rightarrow$  2 CPUs  $\Rightarrow$  150 each in parallel

⑥ Time to finally apply cache affinity  $\rightarrow$  -A

-n 2 -L a:100:100, b:100:50, c:100:50  
-A a:0, b:1, c:1

CPU 0  $\Rightarrow$  runs a

CPU 1  $\Rightarrow$  runs b, c

CPU 0  
Job a 0-9  
cache warmed for Job a  
using 100 of cache-size  
- 90 Job a

Job a 9-29  
= 70 Job a

= 90 Job c

Job a 29-49

Job b 9-29

CPU 1  
Job b 0-9  
Cache warmed for Job b  
using 50 of cache-size  
- 90 Job b

Job c 0-9  
cache warmed for Job c  
using 50 of cache-size (100)

$$= 50 \text{ Job}_a$$

$$= 70 \text{ Job}_b$$

$$\begin{aligned} & \text{Job}_a 49-69 \\ & = 30 \text{ Job}_a \end{aligned}$$

$$\begin{aligned} & \text{Job}_c 9-29 \\ & = 70 \text{ Job}_c \end{aligned}$$

$$\begin{aligned} & \text{Job}_a 69-89 \\ & = 10 \text{ Job}_a \end{aligned}$$

$$\begin{aligned} & \text{Job}_b 29-49 \\ & = 50 \text{ Job}_b \end{aligned}$$

$$\begin{aligned} & \text{Job}_a 89-99 \\ & = 0 \text{ Job}_a \end{aligned}$$

$$\begin{aligned} & \text{Job}_c 29-49 \\ & = 50 \text{ Job}_c \end{aligned}$$

$$\begin{aligned} & \text{Job}_b 49-69 \\ & = 30 \text{ Job}_b \end{aligned}$$

$$\begin{aligned} & \text{Job}_c 49-69 \\ & = 30 \text{ Job}_c \end{aligned}$$

$$\begin{aligned} & \text{Job}_b 69-89 \\ & = 70 \text{ Job}_b \end{aligned}$$

$$\begin{aligned} & \text{Job}_c 89-99 \\ & = 10 \text{ Job}_c \end{aligned}$$

Job b 89-99

= 0 Job b

Job c 89 89

= 0 Job c

Finished time 110

Per-CPU stats

CPU 0 utilization 50.00 [ warm 40.91 ]  
CPU 1 utilization 100.00 [ warm 81.82 ]

Other combinations of 'a', 'b' and 'c'  
will run slower since b+c  
uses cache only because  $50+50 =$   
cache-size in all other cases  
with 'a' it will overwritten.

⑦ Super-linear speedup  $\Rightarrow$  When you run on N CPUs and get a speedup of more than a factor of N.

-L a:100:100, b:100:100, c:100:100 -M 50

Run of  $n=1, 2, 3$   
 $\downarrow$   
CPUs

$n=1$ , -M 50

$$\text{Total time} = 100 + 100 + 100 = 300$$

since cache-size < working set size  
for all jobs

$n=2$ , -M 50

$$\text{Total time} = \frac{100+100+100}{2} = \frac{300}{2} = 150$$

$n=3$ , -M 50

$$\text{Total time} = \frac{100+100+100}{3} = \frac{300}{3} = 100$$

Now let's try with  $M=100$

$n=1$ .

$$\text{Total time} = 100 + 100 + 100 = 300$$

This time cache can be warmed but get's overwritten by next job run

$$n=2$$

$$\text{Total time} = \underbrace{100 + 100 + 100}_n = \frac{300}{2} = 150$$

Again cache overwritten

$$n=3$$

$$\text{Total time} = \underbrace{(10+45) + (10+45) + (10+45)}_n$$

$$= \frac{165}{3} = \underline{\underline{55}}$$

$$\text{warm} = \frac{45}{55} = 81.82\%$$

Each CPU has its own job  
hence cache becomes effective

$$\text{Speed up now is } \frac{300}{55} = 5.46 \times \frac{N}{\text{CPUs}}$$

Super-linear speed up  
Now I understand the name

I guess since each CPU has a specific job running leading to further speed-up due to cache.

```
Finished time 55  
Per-CPU stats  
CPU 0 utilization 100.00 [ warm 81.82 ]  
CPU 1 utilization 100.00 [ warm 81.82 ]  
CPU 2 utilization 100.00 [ warm 81.82 ]
```

⑧ per-CPU scheduling option, -P flag

-L a:100:100, b:100:50, c:100:50

with cache affinity before (⑥)  
Total time = 110

With per-CPU scheduling  
 $-P \Rightarrow$  how often to peek at other schedule queue.  
Default = 30

CPU 0

$Q_0 \rightarrow A \rightarrow C$

CPU 1

$Q_1 \rightarrow B$

1

$$(10 + 45) = 55$$

30 (CPU not free)

$$30 + 10 + 10 + 10 = 60$$

Peak (CPU free)

60 (CPU free,  
takes "a" job)

70 of c remaining  
using cache

$$+ 35 = 95$$

$$\frac{95}{100}$$

70 of a remaining  
10 run "a"

using cache

$$60/2 = 30$$

$$60 + 10 + 30 = 100$$

$$\frac{95}{100}$$

Finished time 100

Per-CPU stats

CPU 0 utilization 95.00 [ warm 35.00 ]  
CPU 1 utilization 95.00 [ warm 75.00 ]

since free during  
(55-60)

With  $-P = 10, 20$  no change since no  
idle job to steal, the other one just ran

with  $-P = 5$ ,

$$Q_0 = a \rightarrow c , Q_1 = b$$

CPU 0

CPU 1

On five seconds  
sees Q<sub>0</sub> and  
steals "C" since  
idle

Stealing happens to and fro  
but CPU 1 steals after hence  
always gets the additional job

Once each queue has single job  
then no stealing since no idle job.

With -P reduced to 5, the  
time is 90.

```
Finished time 90
Per-CPU stats
CPU 0 utilization 94.44 [ warm 72.22 ]
CPU 1 utilization 94.44 [ warm 72.22 ]
```

With -P higher like 40, we  
get more time

```
Finished time 115
Per-CPU stats
CPU 0 utilization 95.65 [ warm 26.09 ]
CPU 1 utilization 78.26 [ warm 60.87 ]
```

With CPUs increase to 3 , best time SS is achieved with 100% CPU utilization and fairness since each job gets a CPU, not idle .

CPU beyond the number of jobs don't help in perf getting any better Per-CPU approach especially good for fairness among jobs and balancing among CPUs.

- ⑨ Played around with various options combinations.